

Practical 4

Title: Data Analytics I

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")
```

```
[ ]: column_names = ['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS',
    ↳ 'RAD', 'TAX', 'PTRATIO', 'B', 'LSTAT', 'PRICE']
data = pd.read_csv('housing.csv', header=None, delimiter=r"\s+",
    ↳ names=column_names)
data.head()
```

```
[ ]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	

	PTRATIO	B	LSTAT	PRICE
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2

```
[ ]: data.isnull().sum()
```

```
[ ]:
```

CRIM	0
ZN	0
INDUS	0
CHAS	0
NOX	0
RM	0
AGE	0
DIS	0

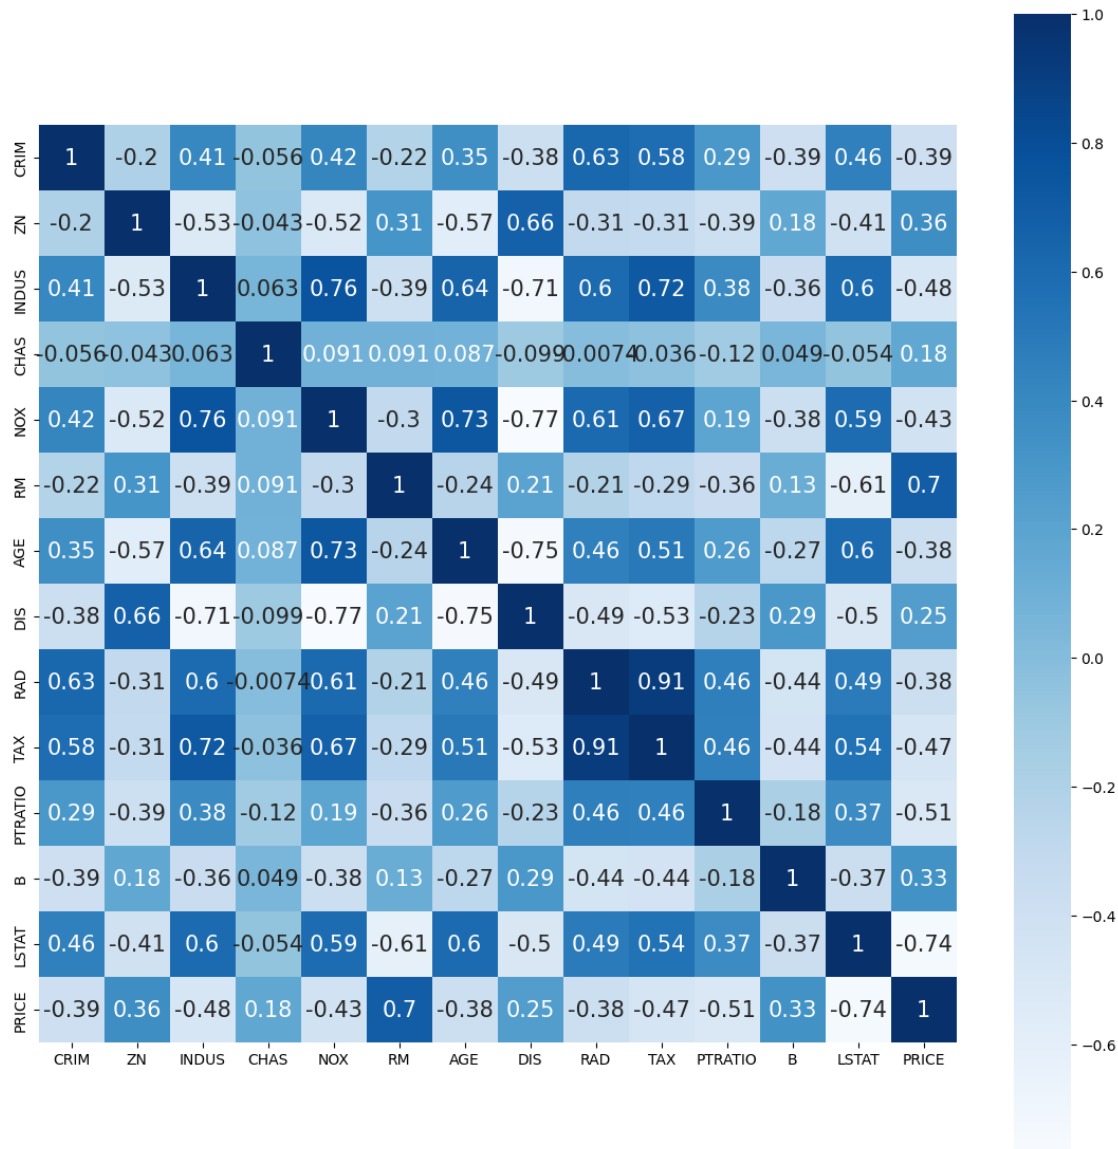
```
RAD      0
TAX      0
PTRATIO  0
B        0
LSTAT    0
PRICE    0
dtype: int64
```

```
[ ]: corr = data.corr()
     corr.shape
```

```
[ ]: (14, 14)
```

```
[ ]: plt.figure(figsize=(14,14))
     sns.heatmap(corr, cbar=True, square=True, annot=True, annot_kws={'size':15},
     ↪ cmap='Blues')

     plt.show()
```



```
[ ]: x = data.drop(['PRICE'], axis = 1)
      y = data['PRICE']
```

```
[ ]: from sklearn.model_selection import train_test_split
      xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size = 0.
      ↪ 2, random_state = 0)
```

```
[ ]: import sklearn
      from sklearn.linear_model import LinearRegression
      lm = LinearRegression()
      # Train the model using the training sets
      model = lm.fit(xtrain, ytrain)
```

```
[ ]: xtrain
```

```
[ ]:      CRIM      ZN  INDUS  CHAS    NOX     RM   AGE     DIS  RAD    TAX  \
220  0.35809   0.0   6.20    1  0.507   6.951   88.5   2.8617    8  307.0
71   0.15876   0.0  10.81    0  0.413   5.961   17.5   5.2873    4  305.0
240  0.11329  30.0   4.93    0  0.428   6.897   54.3   6.3361    6  300.0
6    0.08829  12.5   7.87    0  0.524   6.012   66.6   5.5605    5  311.0
417  25.94060   0.0  18.10    0  0.679   5.304   89.1   1.6475   24  666.0
..    ...    ...    ...    ...    ...    ...    ...    ...
323  0.28392   0.0   7.38    0  0.493   5.708   74.3   4.7211    5  287.0
192  0.08664  45.0   3.44    0  0.437   7.178   26.3   6.4798    5  398.0
117  0.15098   0.0  10.01    0  0.547   6.021   82.6   2.7474    6  432.0
47   0.22927   0.0   6.91    0  0.448   6.030   85.5   5.6894    3  233.0
172  0.13914   0.0   4.05    0  0.510   5.572   88.5   2.5961    5  296.0

      PTRATIO      B  LSTAT
220      17.4  391.70   9.71
71      19.2  376.94   9.88
240      16.6  391.25  11.38
6       15.2  395.60  12.43
417      20.2  127.36  26.64
..      ...    ...    ...
323      19.6  391.13  11.74
192      15.2  390.49   2.87
117      17.8  394.51  10.30
47       17.9  392.74  18.80
172      16.6  396.90  14.69
```

[404 rows x 13 columns]

```
[ ]: ytrain_pred=lm.predict(xtrain)
     ytest_pred=lm.predict(xtest)
```

```
[ ]: testdata=[[0.00632,18.0,2.31,0.0,0.538,6.575,65.2,4.0900,1.0,296.0,15.3,396.
↪90,4.98]]
```

```
[ ]: test_pred = lm.predict(testdata)
     test_pred
```

```
[ ]: array([30.49949836])
```

```
[ ]: df1=pd.DataFrame(ytrain_pred,ytrain)
     df2=pd.DataFrame(ytest_pred,ytest)
     df1
```

```
[ ]:      0
PRICE
26.7    32.556927
```

```

21.7    21.927095
22.0    27.543826
22.9    23.603188
10.4     6.571910
...
18.5    19.494951
36.4    33.326364
19.2    23.796208
16.6    18.458353
23.1    23.249181

```

[404 rows x 1 columns]

```

[ ]: from sklearn.metrics import mean_squared_error, r2_score
mse = mean_squared_error(ytest, ytest_pred)
print('MSE on test data:',mse)
mse1 = mean_squared_error(ytrain_pred, ytrain)
print('MSE on training data:',mse1)

```

MSE on test data: 33.44897999767657

MSE on training data: 19.326470203585725

```

[ ]: r2 = lm.score(xtest, ytest)
rmse = (np.sqrt(mean_squared_error(ytest, ytest_pred)))
print('r-squared: {}'.format(r2))
print('-----')
print('root mean squared error: {}'.format(rmse))

```


r-squared: 0.5892223849182503

root mean squared error: 5.783509315085138

```

[ ]: #plotting the linear regression model
plt.scatter(ytrain ,ytrain_pred,c='blue',marker='o',label='Training data')
plt.scatter(ytest,ytest_pred ,c='lightgreen',marker='s',label='Test data')
plt.xlabel('True values')
plt.ylabel('Predicted')
plt.title("True value vs Predicted value")
plt.legend(loc= 'upper left')
plt.plot()
plt.show()

```



```
[ ]: testdata=[[0.00632,18.0,2.31,0.0,0.538,6.575,65.2,4.0900,1.0,296.0,15.3,396.
↪90,4.98]]
```

```
[ ]: test_pred = lm.predict(testdata)
test_pred
```

```
[ ]: array([30.49949836])
```